

Tanya Goyal

Email: tanyagoyal@cornell.edu

Website: <https://tagoyal.github.io/>

Academic Positions

Assistant Professor, Cornell University
Department of Computer Science

Fall 2024 -

Postdoctoral Scholar, Princeton University
Princeton Language and Intelligence + Department of Computer Science

2023-2024

Education

University of Texas at Austin
PhD in Computer Science (Advisor: Greg Durrett)
MS in Computer Science
Awarded UTCS's Bert Kay Dissertation Award 2023.

Austin, Texas

2019 – 2023

2017 – 2022

Indian Institute of Technology, Guwahati
B.Tech in Mathematics and Computing

Guwahati, India

2011 – 2015

Select Publications

Challenges in Trustworthy Human Evaluation of Chatbots
Wenting Zhao, Alexander Rush, **Tanya Goyal**
NAACL Findings 2024

One Thousand and One Pairs: A "novel" challenge for long-context language models
Marzena Karpinska, Katherine Thai, Kyle Lo, **Tanya Goyal**, Mohit Iyyer.
EMNLP 2024

LitSearch: A Retrieval Benchmark for Scientific Literature Search
Anirudh Ajith, Mengzhou Xia, Alexis Chevalier, **Tanya Goyal**, Danqi Chen, Tianyu Gao
EMNLP 2024

D2PO: Discriminator-Guided DPO with Response Evaluation Models
Prasann Singhal, Nathan Lambert, Scott Niekum, **Tanya Goyal**, Greg Durrett
COLM 2024

FABLES: Evaluating faithfulness and content selection in book-length summarization
Yekyung Kim, Yapei Chang, Marzena Karpinska, Aparna Garimella, Varun Manjunatha, Kyle Lo, **Tanya Goyal**, Mohit Iyyer
COLM 2024

A Long Way To Go: Investigating Length Correlations in RLHF
Prasann Singhal, **Tanya Goyal**, Jiacheng Xu, Greg Durrett.
COLM 2024

BooookScore: A systematic exploration of book-length summarization in the era of LLMs
Yapei Chang, Kyle Lo, **Tanya Goyal**, Mohit Iyyer.
ICLR 2024

Evaluating Large Language Models at Evaluating Instruction Following
Zhiyuan Zeng, Jiatong Yu, Tianyu Gao, Yu Meng, **Tanya Goyal**, Danqi Chen.
ICLR 2024

WICE: Real-World Entailment for Claims in Wikipedia

Ryo Kamoi, **Tanya Goyal**, Juan Diego Rodriguez, and Greg Durrett.
EMNLP 2023

Understanding Factual Errors in Summarization: Errors, Summarizers, Datasets, Error Detectors

Liyan Tang, **Tanya Goyal**, Alexander Fabbri, Philippe Laban, Jiacheng Xu, Semih Yavuz, Wojciech Kryściński, Justin F. Rousseau, Greg Durrett.
ACL 2023

Shortcomings of Question Answering Based Factuality Frameworks for Error Localization

Ryo Kamoi, **Tanya Goyal**, Greg Durrett.
EACL 2023

NL-Augmenter: A framework for task-sensitive natural language augmentation

Kaustubh D Dhole, et al. (including **Tanya Goyal**)
NEJLT 2023

SNaC: Coherence Error Detection for Narrative Summarization

Tanya Goyal, Jessy Li, Greg Durrett.
EMNLP 2022

HYDRASUM: Disentangling Style Features in Text Summarization with Multi-Decoder Models

Tanya Goyal, Nazneen Rajani, Wenhao Liu, Wojciech Kryściński.
EMNLP 2022

FALTE: A Toolkit for Fine-grained Annotation for Long Text Evaluation

Tanya Goyal, Jessy Li, Greg Durrett.
EMNLP-Demo 2022

Training and Fine-tuning Dynamics for Summarization

Tanya Goyal, Jiacheng Xu, Jessy Li, Greg Durrett.
ACL Findings 2022

Annotating and Modeling Fine-grained Factuality in Summarization

Tanya Goyal, Greg Durrett.
NAACL 2021.

(Workshop) Contemporary NLP Modeling in Six Comprehensive Programming Assignments

Greg Durrett, Jifan Chen, Shrey Desai, **Tanya Goyal**, Lucas Kabel, Yasumasa Onoe, Jiacheng Xu.
Workshop on Teaching NLP 2021

Evaluating Factuality in Generation with Dependency-level Entailment

Tanya Goyal, Greg Durrett.
EMNLP Findings 2020.

Neural Syntactic Preordering for Controlled Paraphrase Generation

Tanya Goyal, Greg Durrett.
ACL 2020.

Embedding time expressions for deep temporal ordering models

Tanya Goyal, Greg Durrett.
ACL 2019.

Your behavior signals your reliability: Modeling crowd behavioral traces to ensure quality relevance annotations

Tanya Goyal, Tyler McDonnell, Mucahid Kutlu, Tamer Elsayed, Matthew Lease.
HCOMP 2018.

An Empirical Analysis of Edit Importance between Document Versions
Tanya Goyal, Sachin Kelkar, Manas Agarwal, Jeenu Grover.
EMNLP 2017

Preprints

Recycled Attention: Efficient Inference for Long-context Language Models
Fangyuan Xu, **Tanya Goyal**^{*}, Eunsol Choi^{*}
arxiv 2024

WildHallucinations: Evaluating Long-form Factuality in LLMs with Real-World Entity Queries
Wenting Zhao, **Tanya Goyal**, Yu Ying Chiu, Liwei Jiang, Benjamin Newman, Abhilasha Ravichander, Khyathi Chandu, Ronan Le Bras, Claire Cardie, Yuntian Deng, Yejin Choi
arxiv 2024

News Summarization and Evaluation in the Era of GPT-3
Tanya Goyal, Jessy Li, Greg Durrett.
arxiv 2022

Previous Work Experience

Salesforce Research Summer 2021
Research Intern
Host: Wojciech Kryściński, worked on text summarization. **Filed 1 US patent based on this work.**

A9 Labs, Amazon Summer 2018
Research Intern
Project: Predictive modelling for customer engagement

Adobe Research, India May 2015 - July 2017
Member of Technical Staff
Prototyped predictive models for e-commerce objectives for Adobe Document Cloud
Granted 12 US patents based on my work here.

Service

Area Chair
ICLR(2025), COLM (2024), LREC-COLING (summarization and generation track, 2024), ARR (multiple cycles, 2023-2024), EMNLP 2023 (summarization track), EACL 2023 (summarization track)

Reviewer (*outstanding reviewer)
EMNLP (2020*, 2021, 2022), CoNLL (2020), NAACL (2021, 2022), ACL (2021, 2022), AKBC (2021), ARR (multiple cycles, 2021-present), TACL (standing reviewer, 2023-present)

Committee Member
ACM India Dissertation Award (2024)

Invited Talks

Semantic Machines Nov 2024
Title: Challenges in Collecting and Learning from Human Feedback

CLunch, UPenn NLP Oct 2024
Title: Challenges in Collecting and Learning from Human Feedback

Clinical and Translational Science Center, Weill Cornell Medicine June 2024
Title: Automatic Evaluation of Chatbots in the era of LLMs

Keynote @ Singapore Symposium on NLP (SSNLP 2023) Title: Evaluation in the GPT-3 Era	December 2023
LLM panel @ IEEE International Conference on Healthcare Informatics Title: Evaluating Text Generation in the GPT-3 Era	June 2023
Microsoft Research Title: Building Reliable Text Generation Models	March 2023
NLP Seminar Series at AI Sweden (link) Title: Text summarization and evaluation in the era of GPT-3	December 2022
Adobe Research India Title: Bridging the gap between generation and evaluation capabilities of large language models	November 2022
University of Alberta Title: Syntax-controlled generation using deep neural models	July 2019

Mentorship

• Oliver Li (PhD at Cornell University)	Fall 2024 - Present
• Omer Mustafa Gul (PhD at Cornell University)	Spring 2025 - Present
• Chengyu Huang (PhD at Cornell University)	Fall 2024 - Present
• Marcus Posey (Undergrad at Cornell University)	Fall 2024 - Present
• Vivian Chen (Undergrad at Cornell University)	Fall 2024 - Present
• Rose Xu (Undergrad at Cornell University)	Fall 2024 - Present
• Prasann Singhal (Undergrad at UT Austin)	Fall 2023 - Present
• Ryo Kamoi (MS at UT Austin, next PhD UPitt)	Spring 2022 - Spring 2023
• Liyan Tang (PhD at UT Austin)	Spring 2022
• Sachin Kelkar (Research Intern at Adobe Research, next MS at Columbia)	Summer 2016
• Manas Agarwal (Research Intern at Adobe Research)	Summer 2016
• Jeenu Grover (Research Intern at Adobe Research, next Google India)	Summer 2016

Patents *Granted by US PTO*

Filed by Salesforce Research Multi-attribute control for text summarization using multiple decoder heads	<i>US Patent 11,755,637</i>
Filed by Adobe Research Generating data-driven geo-fences	<i>US Patent 9,838,843</i>
Method to expand seed keywords into a relevant social query	<i>US Patent 9,965,766</i>
Generating predictive models for authoring short messages	<i>US Patent 10,528,652</i>
Content to layout template mapping and transformation	<i>US Patent 10,521,494</i>
Classifying and ranking changes between document revisions	<i>US Patent 10,713,432</i>
Constructing enterprise-specific knowledge graphs	<i>US Patent 10,915,577</i>
Determination of Paywall Metrics	<i>US Patent 10,853,887</i>
Prediction of tone of interpersonal text communications	<i>US Patent 10,796,095</i>
Tagging documents with security policies	<i>US Patent 11,748,501</i>
Content optimization for audiences	<i>US Patent 10,922,492</i>
Bundling online content fragments for presentation based on content-specific metrics and inter-content constraints	<i>US Patent 10,891,667</i>
Method and apparatus for generating predictive insights for authoring short messages	<i>US Patent 10,073,822</i>